STAGE 1: PERCEPTUAL COMPRESSION (VARIATIONAL AUTOENCODER)

KL-regularized latent space via probabilistic encoding and reparameterization

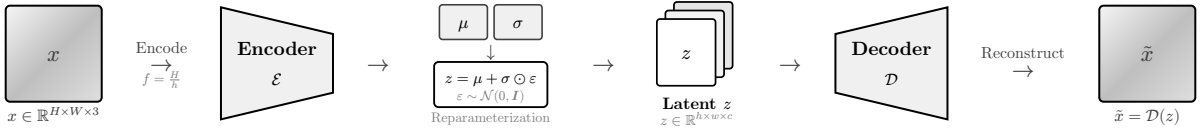

STAGE 2: LATENT DIFFUSION PROCESS
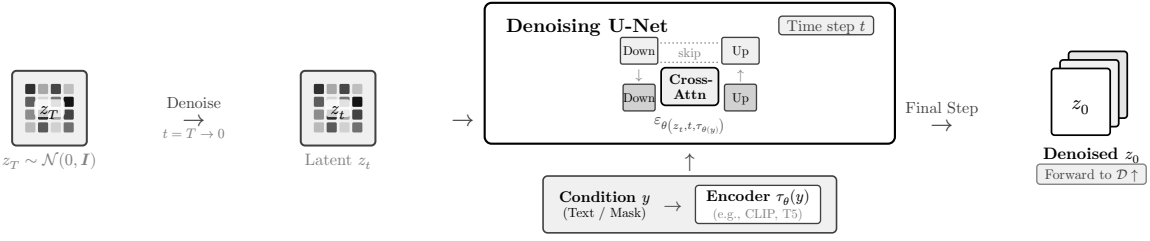
Iterative denoising directly in latent space, conditioned via Cross-Attention