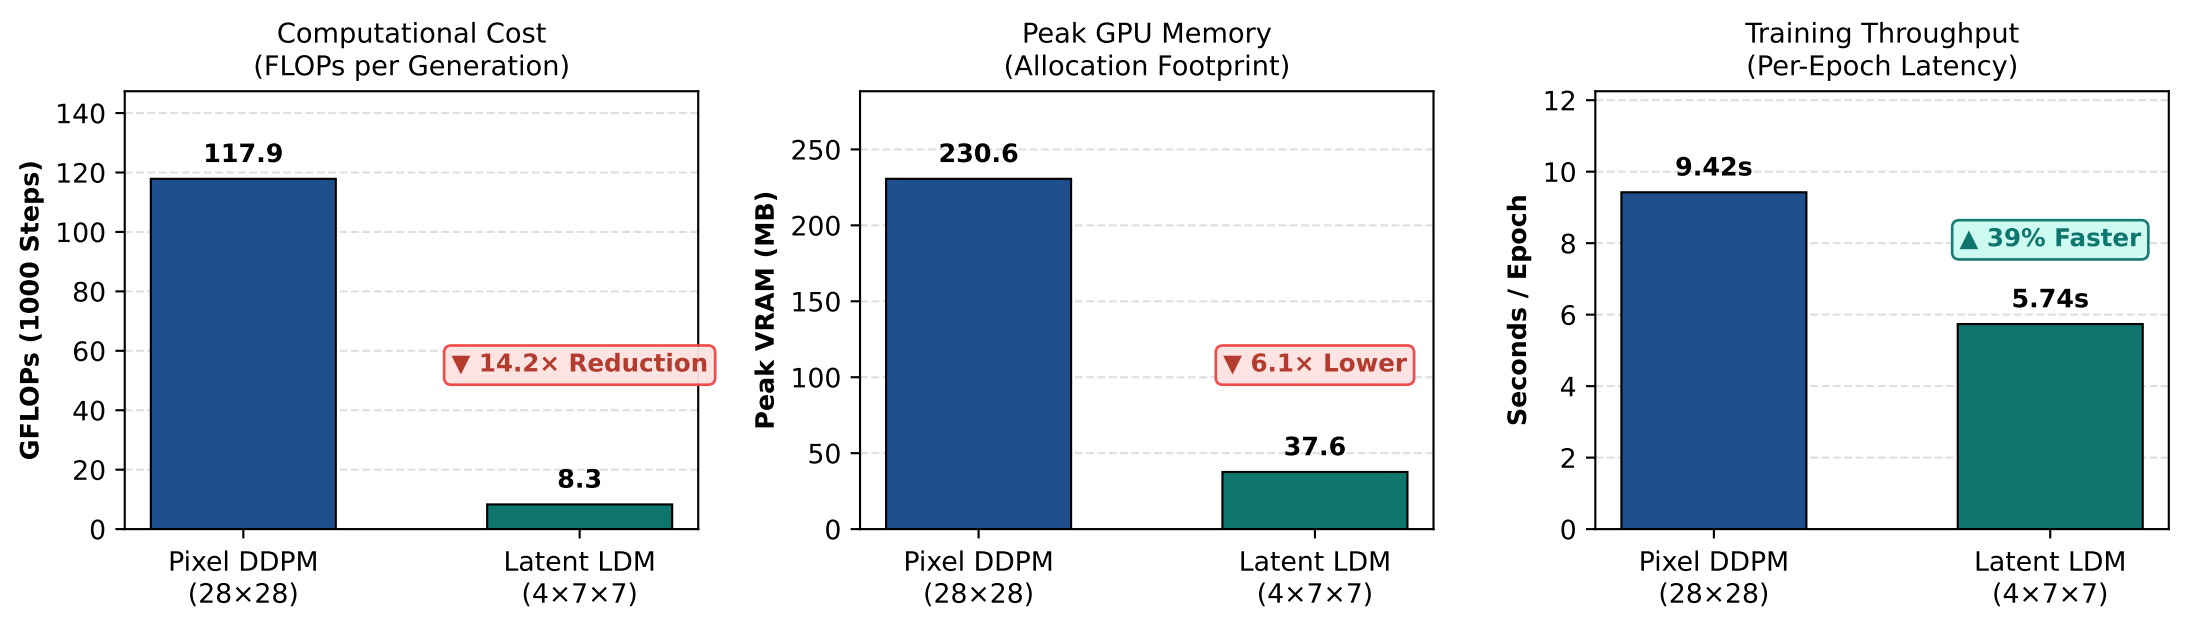
Computational Cost
(FLOPs per Generation)
GFLOPs (1000 Steps)
140
120
100
80
60
40
20
0
117.9
8.3
▼ 14.2× Reduction
Pixel DDPM
(28×28)
Latent LDM
(4×7×7)
Peak GPU Memory
(Allocation Footprint)
Peak VRAM (MB)
250
200
150
100
50
0
230.6
37.6
▼ 6.1× Lower
Pixel DDPM
(28×28)
Latent LDM
(4×7×7)
Training Throughput
(Per-Epoch Latency)
Seconds / Epoch
12
10
8
6
4
2
0
9.42s
5.74s
▲ 39% Faster
Pixel DDPM
(28×28)
Latent LDM
(4×7×7)